



TECHNICAL REPORT

Crackathon – IIT Bombay
Road Damage Detection using Deep Learning

Github Link : <https://github.com/VishalKR-12/Crackathon>

1. Introduction

Road infrastructure plays a vital role in transportation safety and economic development. However, manual inspection of road surfaces for damages such as cracks and potholes is inefficient, time-consuming, and prone to subjectivity. To address this challenge, this project presents an **automated road damage detection system** using deep learning-based object detection.

The objective of this work is to develop a robust computer vision model that can **accurately detect and classify road damages** from images using the **RDD2022 dataset**, following the YOLO detection format. The system aims to achieve high precision and recall under real-world conditions and perform well on a hidden test dataset evaluated using **Mean Average Precision (mAP)**.

2. Dataset Description

The project uses the **Road Damage Detection 2022 (RDD2022)** dataset, provided by the Crackathon organizers.

Dataset Characteristics:

- **Total Images:** ~47,000 high-resolution road images
- **Data Split:**
 - Training set (with labels)
 - Validation set (with labels)
 - Test set (images only, labels hidden)
- **Label Format:** YOLO TXT format
- **Damage Classes:**

Class ID Damage Type

0	Longitudinal Crack
1	Transverse Crack
2	Alligator Crack
3	Other Corruption
4	Pothole

To ensure fair evaluation, **only the randomized dataset provided by the organizers** was used, and no external data was included.

3. Model Architecture

The object detection model is based on **YOLOv8**, a state-of-the-art single-stage detector known for its balance between speed and accuracy.

Architecture Highlights:

- Backbone: CSP-based convolutional layers
- Neck: Feature Pyramid Network (FPN + PAN)
- Head: Anchor-free detection head
- Loss Functions:
 - Bounding box regression loss
 - Classification loss
 - Objectness confidence loss

A pre-trained **YOLOv8-Medium (YOLOv8m)** model was used as the starting point to leverage transfer learning, improving convergence speed and detection accuracy.

4. Training Strategy

Fine-Tuning

- The model was fine-tuned on the training split of the RDD2022 dataset.
- Pretrained weights were used to improve generalization.
- Training was conducted using Ultralytics YOLO framework.

Key Training Parameters:

- Image size: 640×640
 - Optimizer: Default YOLO optimizer (AdamW)
 - Batch size: Hardware-dependent (Kaggle GPU)
 - Training epochs: 18
 - Confidence threshold: 0.10
 - Evaluation metric: mAP@0.50–0.95
-

5. Data Augmentation

To improve robustness against varying lighting conditions, road textures, and camera angles, multiple augmentation techniques were applied during training:

- Random horizontal flipping
- Scaling and resizing
- Mosaic augmentation
- Color jittering
- Random cropping

These augmentations helped reduce overfitting and improved the model's ability to generalize to unseen test images.

6. Hyperparameter Tuning and Optimization

Multiple controlled experiments were conducted to improve model performance over the baseline:

- Comparison between **partial layer freezing** and **full fine-tuning**
- Adjustment of confidence thresholds during inference
- Batch size experimentation within GPU memory limits
- Evaluation of **Test-Time Augmentation (TTA)**

Model improvements were selected based on validation mAP scores, with full fine-tuning and TTA providing the most consistent gains.

7. Training Dynamics and Performance Analysis

The YOLOv8m model was trained for **18 epochs**, during which both training and validation metrics showed **consistent and stable improvement**, indicating good convergence without overfitting.

Loss Convergence

Across epochs, all three major loss components showed a steady downward trend:

- **Box Loss:** Reduced from **1.75 → 1.47**
- **Classification Loss:** Reduced from **2.09 → 1.39**
- **DFL (Distribution Focal Loss):** Reduced from **1.59 → 1.43**

This reduction confirms improved localization accuracy, better class separation, and more confident bounding box predictions.

Validation losses closely tracked training losses, indicating **strong generalization** and **minimal overfitting**.

8. Quantitative Evaluation Results

Model performance was evaluated using standard object detection metrics on the validation set.

Final Metrics (Epoch 18 – Best Model)

Metric	Value
Precision	0.608
Recall	0.551
mAP@0.50	0.584
mAP@0.50–0.95	0.311

Observations

- **Precision improvement** from **0.41 → 0.61** indicates fewer false positives.

- Recall improvement from **0.35 → 0.55** indicates better detection of road damage instances.
- mAP@0.50 steadily increased from **0.31 → 0.58**, showing strong detection accuracy.
- mAP@0.50–0.95, which is stricter and more challenging, reached **0.31**, a solid result for multi-class road damage detection.

These trends confirm that the model learned meaningful visual representations of different damage types.

9. Learning Rate Scheduling

A cosine-based learning rate decay strategy was used during training.

- Initial learning rate: ~3.7e-4
- Final learning rate: ~7.2e-5

Gradual decay helped:

- Stabilize late-stage training
- Prevent oscillations
- Improve final mAP scores

This scheduling contributed to smooth convergence and consistent validation improvements.

10. Impact of Test-Time Augmentation (TTA)

While validation metrics were computed without TTA, inference on the test set used **Test-Time Augmentation**, which is known to:

- Improve recall for small cracks
- Increase robustness to orientation and lighting changes
- Reduce missed detections

Due to memory constraints, TTA inference was performed using a **chunk-based strategy (200 images per batch)**, ensuring safe execution in the Kaggle environment without kernel crashes.

11. Post-Processing

To ensure that **every test image has a corresponding prediction file**, a guaranteed post-processing step was implemented.

Post-Processing Steps:

- Iterate through all test images
- If prediction exists → filter detections using confidence threshold
- If no detection exists → create an empty .txt file
- Ensures **100% submission compliance**

Each output file follows the required format:

```
<class_id> <x_center> <y_center> <width> <height> <confidence_score>
```

12. Final Submission Packaging

- All prediction .txt files were placed in a single folder
- The folder was compressed into submission.zip
- The trained model (best.pt) and source code were packaged separately

This structure strictly follows Crackathon submission guidelines.

13. Conclusion

This project successfully demonstrates a scalable and automated solution for road damage detection using deep learning. By combining fine-tuned YOLOv8 architecture, extensive data augmentation, test-time augmentation, and robust post-processing, the system achieves strong performance while remaining practical for real-world deployment.

The final model achieved a **validation mAP@0.50 of 58.4%** and **mAP@0.50–0.95 of 31.1%**, demonstrating strong detection capability across diverse road damage categories. The consistent improvement across all metrics confirms the effectiveness of the training strategy, augmentation techniques, and architectural choices.

This solution provides a scalable and practical AI-based approach for automated road damage detection, suitable for real-world deployment scenarios.



